A decorative graphic on the left side of the slide, consisting of a network of thin, light-blue lines and small circles, resembling a circuit board or a neural network, extending vertically from the top to the bottom.

# COSC 416: Topics in Databases (DBaaS)

TOPIC 2: DBAAS ARCHITECTURE

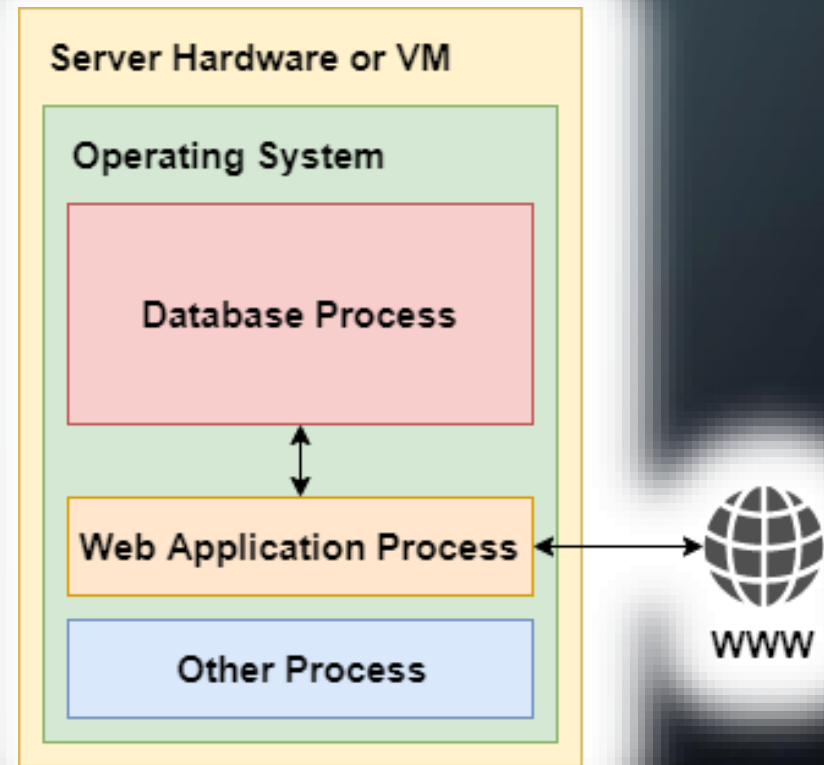
# SCHEDULE

- 1. DBaaS Architecture Rundown*
- 2. Benefits and Drawbacks of DBaaS Architecture*
- 3. A look at Amazon RDS*

# DBAAS ARCHITECTURE

# A TRADITIONAL DATABASE SERVER

- Let's consider a traditional database server for a moment
- The server may be a standalone database server or shared with an application
- The server may be on physical hardware or on a virtual machine



# A HOLISTIC APPROACH

- Using actual hardware or a VM, we are dealing with an entire *software stack* (and potentially hardware stack) to run our database
- This software stack includes the operating system, the DBMS, and any dependencies for the DBMS
- System configuration is ultimately up to the system administrator or developer

# WHAT ABOUT CLUSTERS?

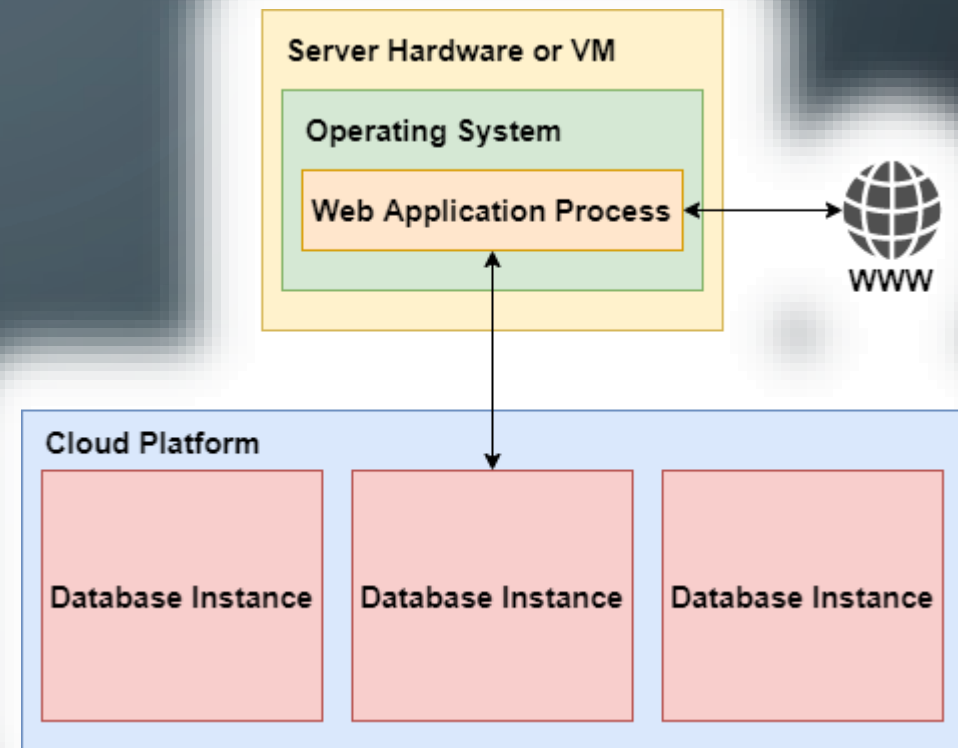
- You might wonder how cluster-based databases like MongoDB fit into this
- Traditional installations still require that each node be configured → Still using a software stack
- Multiplies work, but also allows for things like replication (automatic recovery) and better scaling through the use of multiple nodes

## SO HOW IS DBAAS DIFFERENT?

- Using a Database-as-a-Service gives us a database process, without the underlying *software stack*
- This eliminates the need to handle the operating system, dependencies, and underlying server itself

# DBAAS ARCHITECTURE

- Instead of running the database ourselves, we offload that work to a cloud platform
- The database instance is accessed remotely by the client or application process





# WHAT EXACTLY IS “THE CLOUD”

- Amazon likes to helpfully point out that RDS is hosted “In the cloud”, but what does that mean exactly?
- The cloud is really just a collection of networked dedicated servers, which run hypervisors that manage individual guest machines, which in turn run the database instances
- The guest machines are preconfigured for the database and since they are virtual, they are disconnected from the hardware itself (aiding scalability)

# HOW DOES DBAAS DIFFER FROM A VM?

- DBaaS systems are very much like virtual machines that you might rent from Digital Ocean or Linode
- Where they differ is the level of *management* offered
- With a virtual machine, you're still usually responsible for installing and maintaining your software stack
- Additionally, most VMs aren't specialized for database use (emphasis on network connection speed and disk I/O speed)

# MANAGED SOFTWARE SOLUTIONS

- The idea of a *managed software solution* is that it is easy and fast to create, use, and teardown a database instance, without having to deal with the messy parts (configuration and maintenance)
- Virtual machines remove the need for hardware management in servers, while DBaaS solutions remove the need for software management, beyond the DBMS itself

# AT THE ENTERPRISE LEVEL

- At a large-scale enterprise level, managed software solutions may even include regular database maintenance and design
- This includes having dedicated database engineers and administrators on hand to help design and maintain the database itself
- Totally managed solutions are, understandably, quite expensive, but are considered worthwhile at the enterprise level, where downtime could result in a considerable cost

# COMMUNICATING WITH THE DATABASE

- In a traditional database server architecture, we interact with the database in one of two ways:
  - Inside the database console, performing CRUD operations either locally or via remote connection
  - Outside the database console: creating backups, sizing databases, starting and stopping the DBMS

# COMMUNICATING WITH A DBAAS

- With a DBaaS system, we can still connect to the database directly to perform database operations via a remote connection
- Since we usually don't have access to the Operation System (Often no SSH access), how do we perform administrative tasks?
  - Most DBaaS solutions offer an API for administrative tasks

# DBAAS API CONTROL

- The DBaaS API will usually allow you to create, modify, destroy, stop, start, and resize database instances
- This is done via a standard HTTPS web API, allowing your application to use standard web requests to manage your DBMS instances
- There's a security element to this (requests require a secret key), so not just anyone can make these requests



# A CAVEAT ABOUT THE TERMINOLOGY

- As with most things at the forefront of technology, the term DBaaS is used pretty loosely (see: “the cloud”)
- Some services describe themselves as DBaaS when they’re more of a preconfigured, but still unmanaged virtual machine
- Some DBaaS services don’t offer APIs, and simply offer standalone database instances without administrative functions
- The general idea however, is that *DBaaS systems use a preconfigured instance optimized for databases without requiring configuration by the client.*



The background is a dark blue gradient. In the corners, there are white line-art illustrations of circuit boards or neural networks, with lines connecting to small circles.

# PROS AND CONS OF DBAAS SYSTEMS

# BENEFITS OF TRADITIONAL DATABASE SERVERS

- There are some benefits to using a traditional database architecture:
  - You have complete control over the system – operating system, dependencies, etc
  - If you already have a server or VM for your application, it's often pretty trivial to spin up a database on the server

# LONG STORY SHORT

- For small applications, or applications where the database size isn't going to fluctuate greatly and is small in size, traditional database systems are fine
- There's not necessarily a point to making things more complicated than needed if you just want a MySQL database for your application

# DRAWBACKS TO A TRADITIONAL ARCHITECTURE

- There are a lot of drawbacks though:
  - Hardware or VM rental costs (\$)
  - Software licensing costs (Paid enterprise-level DBMS products)
  - Requirement that you perform maintenance and patches as required
  - Requires monitoring to maintain uptime, as well as a good stable host server or VM
  - Depending on the configuration, can become complicated quickly (clusters, running multiple databases at once on the same machine, handling remote access securely)

# WHERE DBAAS SHINES

- DBaaS systems really shine in providing enterprise-level databases at a fraction of the cost
  - Fees based on instance size and usage
  - No need to pay for hardware
  - Software licensing fees usually not required, or wrapped into usage fees
  - Is optimized for database usage and doesn't require regular maintenance or upkeep, is monitored by the provider

# DBAAS SCALABILITY

- DBaaS systems are also very preferable if we need a *scalable* system
- Database instances can be quickly resized to accommodate larger or smaller databases
- This is generally a lot smoother than trying to resize a physical server, or resizing a standalone virtual machine (inherent OS issues, and potentially requiring that the database be taken down during resizing)

# ADDITIONAL VALUE-ADDED BENEFITS

- Depending on the DBaaS you use, there may be additional benefits:
  - Pre-baked backup systems that allow you to take database snapshots and perform recoveries quickly
  - HTTP APIs that simplify application management of the database
  - Access to many different database instance types, with a common API (for example, quickly spinning up MySQL, Oracle, and PostgreSQL instances via HTTP)

# DRAWBACKS TO DBAAS SYSTEMS

- DBaaS systems are of course, not perfect:
  - Some control is ceded to the database provider; you lack access to the operating system and dependencies
  - Problems on the provider's end can jeopardize your database (both in terms of uptime, and security)
  - Costs can rack up quickly if you have to resize to larger instances



# SOMETHING ELSE TO CONSIDER: DATA LAWS

- With more legislation being passed to protect data privacy and security, DBaaS solutions may not be the best way to handle some data
- Sensitive data like medical records may be legally required to be stored in-country
  - Can you guarantee your DBaaS instance will be hosted in Canada?
- Putting such data on cloud servers would be a massive liability – in these cases, local datacenters or private servers would be a much better option for safeguarding the data

# SO, WHEN TO USE A DBAAS?

- A DBaaS system makes sense if you need a rapidly scalable database system, or the ability to quickly create and destroy database instances
- It also makes sense if you want to minimize the maintenance and monitoring of your database, in particular if you're running many instances or using clusters
- Generally, a good option for businesses that want enterprise database capabilities but can't justify dedicated local hardware
  - Popular with start-ups, where operating capital is constrained and the ability to quickly wind down database servers is valuable if the business fails

# DBAAS FOR OTHER APPLICATIONS

- So far, we've been approaching this from the point of view of a web application
- However, DBaaS services offer great potential for desktop applications
- Allows you to completely eliminate a standalone server, and have your desktop application directly interface with the DBaaS instance
- This also avoids the bottleneck of using a single application–database connection by giving each user their own connection

# AMAZON RDS PREVIEW

# AMAZON RELATION DATABASE SERVICE

- The Amazon Relational Database Service (RDS) is a relational database DBaaS provided by Amazon, as a companion product to their AWS virtual server service
- RDS is a general database platform that offers many common relational database engines

# SUPPORTED DBMS'S

- RDS offers the following supported databases:
  - MySQL
  - MariaDB
  - MS SQL
  - Oracle DB
  - PostgreSQL

# RDS FEATURES

- Amazon RDS offers a scalable database instances with a common HTTP API for management
- Automated backups and replication
- Built-in monitoring
- Support for replication as both a means of backup, and dealing with increased traffic (load can be distributed across replicas)

# RDS SIZING AND PRICING

- Amazon RDS offers instances ranging in size from single-core micro machines with 1 GB of memory, up to 64-core machines with massive memory (500 GB) and substantial throughput
- Pricing is based on a combination of instance size, instance usage, and software license costs depending on the DBMS chosen






## NEXT WEEK

- Next week we'll be diving right in to creating Amazon RDS machines, and communicating with them from an application
- We'll talk about the RDS web interface, as well as the API used to manage our DB instances, and how we can integrate it into our application



# SCHEDULE

- 1. DBaaS Architecture Rundown*
  - 2. Benefits and Drawbacks of DBaaS Architecture*
  - 3. A look at Amazon RDS*
- 
- 
- 

The image features a dark blue gradient background with faint, stylized circuit board traces in the corners. These traces are composed of thin white lines and small white circles, resembling electronic components or data paths. The central text is a large, white, sans-serif font that reads "SO LONG, FOLKS!".

SO LONG, FOLKS!